

Name : Abhang Rushikesh
Roll No : BEA-30

Group B: Machine Learning

Assignment B2

Classify the email using the binary classification method.

Email Spam detection has two states: a) Normal State – Not Spam, b) Abnormal State – Spam.

Use K-Nearest Neighbors and Support Vector Machine for classification. Analyze their performance. Dataset link: The emails.csv dataset on the Kaggle

<https://www.kaggle.com/datasets/balaka18/email-spam-classification-dataset-csv>

<https://www.kaggle.com/datasets/balaka18/email-spam-classification-dataset-csv>

```
In [1]: import numpy as np
import pandas as pd
```

```
In [2]: df = pd.read_csv("emails.csv")
```

```
In [3]: df.head()
```

Out[3]:

Email No.	the	to	ect	and	for	of	a	you	hou	...	connevey	jay	valued	lay	infrastructure
-----------	-----	----	-----	-----	-----	----	---	-----	-----	-----	----------	-----	--------	-----	----------------

0	0	0	1	0	0	0	2	0	0	...	0	0	0		
1	8	13	24	6	6	2	102	1	27	...	0	0	0		
2	0	0	1	0	0	0	8	0	0	...	0	0	0		
3	0	5	22	0	5	1	51	2	10	...	0	0	0		
4	7	6	17	1	5	2	57	0	9	...	0	0	0		

5 rows x 3002 columns

In [4]: `df.columns`

Out[4]: Index(['Email No.', 'the', 'to', 'ect', 'and', 'for', 'of', 'a', 'you', 'hou',
...
 'connevey', 'jay', 'valued', 'lay', 'infrastructure', 'military',
 'allowing', 'ff', 'dry', 'Prediction'],
 dtype='object', length=3002)

In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5172 entries, 0 to 5171
Columns: 3002 entries, Email No. to Prediction
dtypes: int64(3001), object(1) memory usage:
118.5+ MB
```

In [6]: `df.shape`

Out[6]: (5172, 3002)

In [7]: `df.duplicated().sum()`

Out[7]: 0

In [8]: `df.describe()`

Out[8]:

	the	to	ect	and	for	of	a
count	5172.000000	5172.000000	5172.000000	5172.000000	5172.000000	5172.000000	5172.000000
mean	6.640565	6.188128	5.143852	3.075599	3.124710	2.627030	55.517401
std	11.745009	9.534576	14.101142	6.045970	4.680522	6.229845	87.574172
min	0.000000	0.000000	1.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	1.000000	1.000000	0.000000	1.000000	0.000000	12.000000
50%	3.000000	3.000000	1.000000	1.000000	2.000000	1.000000	28.000000
75%	8.000000	7.000000	4.000000	3.000000	4.000000	2.000000	62.250000
max	210.000000	132.000000	344.000000	89.000000	47.000000	77.000000	1898.000000

8 rows x 3001 columns

In

In [9]: `# df.corr()`

In [10]: `from sklearn.model_selection import train_test_split`
`from sklearn.svm import SVC`
`from sklearn.metrics import accuracy_score`

In [11]: `df[df.isnull().any(axis=1)]`

Out[11]:

Email	the	to	ect	and	for	of	a	you	hou	...	connevey	jay	valued	lay	infrastructure	m
No.																

0 rows x 3002 columns

[12]: `df.Prediction.value_counts()`

Out[12]:

	0	3672	1
1500			

Name: Prediction, dtype: int64


In [13]: `train,test= train_test_split(df,test_size=0.3,stratify=df.Prediction)`In [14]: `train.shape`

Out[14]: (3620, 3002)

In [15]: `test.shape`

Out[15]: (1552, 3002)

In [16]: `train.pivot_table(index='Prediction',aggfunc='size').plot(kind='bar',title='Class`Out[16]: `<AxesSubplot:title={'center':'Class distribution'}, xlabel='Prediction'>`

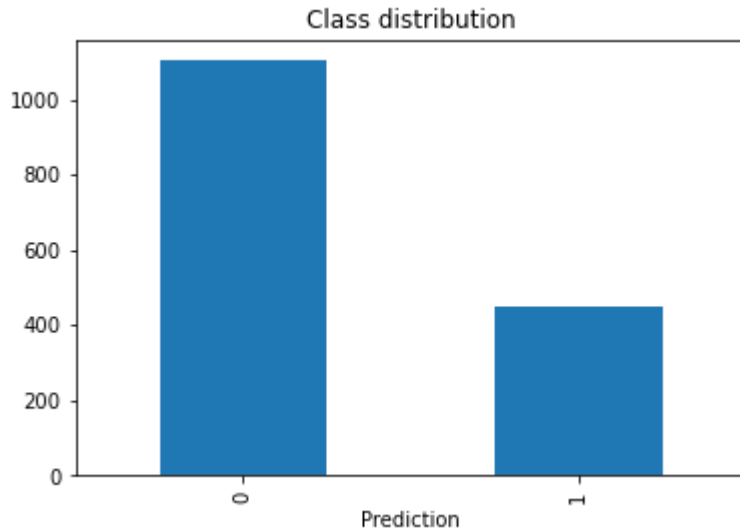


In

[17]:

```
test.pivot_table(index='Prediction',aggfunc='size').plot(kind='bar',title='Class
```

```
Out[17]: <AxesSubplot:title={'center':'Class distribution'}, xlabel='Prediction'>
```



```
In [18]: X = df.iloc[:,1:3001]
```

```
In [19]: Y = df.iloc[:, -1].values
Y
```


```
Out[19]: array([0, 0, 0, ..., 1, 1, 0], dtype=int64)
```

```
In [20]: train_x,test_x,train_y,test_y = train_test_split(X,Y,test_size = 0.25,stratify =
[21]: plt.hist(test_y)
```

```
<IPython.core.display.Javascript object>
```

```
Out[21]: (array([918., 0., 0., 0., 0., 0., 0., 0., 0., 375.]),
array([0. , 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1. ]), <BarContainer
object of 10 artists>)
```

In



SVM

```
In [22]: svc = SVC(C=1.0, kernel='rbf', gamma='auto')
# C here is the regularization parameter. Here, L2 penalty is used(default). It i
# As C increases, model overfits.
# Kernel here is the radial basis function kernel.
# gamma (only used for rbf kernel) : As gamma increases, model
overfits. svc.fit(train_x, train_y) y_pred = svc.predict(test_x)
print("Accuracy Score for SVC : ", accuracy_score(y_pred, test_y))
```

Accuracy Score for SVC : 0.8963650425367363

```
In [23]: acc = accuracy_score(y_pred, test_y)
acc
```

Out[23]: 0.8963650425367363

```
[24]: from sklearn.metrics import precision_score, recall_score, f1_score, fbeta_score, con
```

```
In [25]: pre = precision_score(test_y, y_pred)
pre
```

Out[25]: 0.8801261829652997

In

```
In [26]: recall = recall_score(test_y,y_pred)
recall
```

Out[26]: 0.744

```
In [27]: f1 = f1_score(test_y,y_pred)
f1
```

Out[27]: 0.8063583815028901

```
In [28]: fbeta0_5 = fbeta_score(test_y,y_pred,beta=0.5)
fbeta0_5
```

Out[28]: 0.8490566037735849

```
In [29]: fbeta2 = fbeta_score(test_y,y_pred,beta=2)
fbeta2
```

Out[29]: 0.7677490368739681

```
In [30]: result = pd.DataFrame(columns=['Accuracy score', 'Precision', 'Recall', 'F1 Score',
result.loc['SVM'] = [acc,pre,recall,f1,fbeta0_5,fbeta2] result
```

```
Out[30]:
```

	Accuracy score	Precision	Recall	F1 Score	Fbeta Score(0.5)	Fbeta Score(2)
SVM	0.896365	0.880126	0.744	0.806358	0.849057	0.767749

```
In [31]: confusion_matrix(test_y,y_pred)
```

```
Out[31]: array([[880,  38],
               [ 96, 279]], dtype=int64)
```

```
In [32]: print(classification_report(test_y,y_pred))
```

	precision	recall	f1-score	support
0	0.90	0.96	0.93	918
1	0.88	0.74	0.81	375
accuracy			0.90	1293
macro avg	0.89	0.85	0.87	1293
weighted avg	0.90	0.90	0.89	1293

SMOTE : a powerful solution for imbalanced data

```
In [33]: from imblearn.over_sampling import SMOTE
```

```
In [34]: oversample = SMOTE()
```

```
In [35]: X_sampled, Y_sampled = oversample.fit_resample(X,Y)
```

```
In [36]: X_sampled.shape
```

```
Out[36]: (7344, 3000)
```

```
In [37]: Y_sampled.shape
```

```
Out[37]: (7344,)
```

```
In [38]: X_sampled.head()
```

```
Out[38]:
```

	the	to	ect	and	for	of	a	you	hou	in	...	enhancements	connevey	jay	valued	lay
0	0	0	1	0	0	0	2	0	0	0	...	0	0	0	0	0
1	8	13	24	6	6	2	102	1	27	18	...	0	0	0	0	0
			0	0												
2	0	0	1	0	0	0	8	0	0	4	...	0	0	0	0	0
			0	0												
3	0	5	22	0	5	1	51	2	10	1	...	0	0	0	0	0
			0	0												
4	7	6	17	1	5	2	57	0	9	3	...	0	0	0	0	0
			0	0												

5 rows x 3000 columns


```
In [39]: train_x1,test_x1,train_y1,test_y1 = train_test_split(X_sampled,Y_sampled,test_size=0.2)
```

```
[40]: plt.hist(test_y1)
```

```
<IPython.core.display.Javascript object>
```

```
Out[40]: (array([909., 0., 0., 0., 0., 0., 0., 0., 0., 927.]),
array([0. , 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1. ]), <BarContainer
object of 10 artists>)
```


In



```
In [41]: svc = SVC(C=1.0,kernel='rbf',gamma='auto')
# C here is the regularization parameter. Here, L2 penalty is used(default). It i
# As C increases, model overfits.
# Kernel here is the radial basis function kernel.
# gamma (only used for rbf kernel) : As gamma increases, model overfits.
svc.fit(train_x1,train_y1) y_pred1 = svc.predict(test_x1)
print("Accuracy Score for SVC : ",
      accuracy_score(y_pred1,test_y1))
```

Accuracy Score for SVC : 0.9449891067538126

```
In [42]: confusion_matrix(test_y1,y_pred1)
```

```
Out[42]: array([[842,  67],
                [ 34, 893]], dtype=int64)
```

```
[43]: print(classification_report(test_y1,y_pred1))
```

		precision	recall	f1-score	support
	0	0.96	0.93	0.94	909
1	0.93	0.96	0.95	0.94	927

In

accuracy			0.94	1836
macro avg	0.95	0.94	0.94	1836
weighted avg	0.95	0.94	0.94	1836

```
In [44]: acc = accuracy_score(y_pred1,test_y1)
acc
```

```
Out[44]: 0.9449891067538126
```

```
In [45]: pre = precision_score(test_y1,y_pred1)
pre
```

```
Out[45]: 0.9302083333333333
```

```
In [46]: recall = recall_score(test_y1,y_pred1)
recall
```

```
Out[46]: 0.9633225458468176
```

```
In [47]: f1 = f1_score(test_y1,y_pred1)
f1
```

```
Out[47]: 0.9464758876523581
```

```
In [48]: fbeta0_5 = fbeta_score(test_y1,y_pred1,beta=0.5)
fbeta0_5
```

```
Out[48]: 0.936647786868051
```

```
In [49]: fbeta2 = fbeta_score(test_y1,y_pred1,beta=2)
fbeta2
```

```
Out[49]: 0.9565124250214223
```

```
In [50]: result.loc['SVM_SMOTE'] = [acc,pre,recall,f1,fbeta0_5,fbeta2]
result
```

```
Out[50]:
```

	Accuracy score	Precision	Recall	F1 Score	Fbeta Score(0.5)	Fbeta Score(2)
SVM	0.896365	0.880126	0.744000	0.806358	0.849057	0.767749
SVM_SMOTE	0.944989	0.930208	0.963323	0.946476	0.936648	0.956512

```
[51]: from sklearn.neighbors import KNeighborsClassifier
KNN= KNeighborsClassifier(n_neighbors=3)
# Train the model using the training sets
KNN.fit(train_x,train_y) y_pred = KNN.predict(test_x)
print("Accuracy Score for KNN : ",
accuracy_score(y_pred,test_y))
```

```
Accuracy Score for KNN : 0.8476411446249034
```

In

```
In [52]: acc = accuracy_score(y_pred,test_y)
acc
```

Out[52]: 0.8476411446249034

```
In [53]: pre = precision_score(test_y,y_pred)
pre
```

Out[53]: 0.7099056603773585

```
In [54]: recall = recall_score(test_y,y_pred)
recall
```

Out[54]: 0.8026666666666666

```
In [55]: f1 = f1_score(test_y,y_pred)
f1
```

Out[55]: 0.7534418022528159

```
In [56]: fbeta0_5 = fbeta_score(test_y,y_pred,beta=0.5)
fbeta0_5
```

Out[56]: 0.7267020762916465

```
In [57]: fbeta2 = fbeta_score(test_y,y_pred,beta=2)
fbeta2
```

Out[57]: 0.7822245322245321

```
In [58]: result.loc['KNN'] = [acc,pre,recall,f1,fbeta0_5,fbeta2]
result
```

Out[58]:

	Accuracy score	Precision	Recall	F1 Score	Fbeta Score(0.5)	Fbeta Score(2)
SVM	0.896365	0.880126	0.744000	0.806358	0.849057	0.767749
SVM_SMOTE	0.944989	0.930208	0.963323	0.946476	0.936648	0.956512
KNN	0.847641	0.709906	0.802667	0.753442	0.726702	0.782225

```
[59]: confusion_matrix(test_y,y_pred)
```

Out[59]: array([[795, 123],
[74, 301]], dtype=int64)

```
In [60]: print(classification_report(test_y,y_pred))
```

		precision	recall	f1-score	support
	0	0.91	0.87	0.89	918
1	0.71	0.80	0.75	375	

In

accuracy			0.85	1293
macro avg	0.81	0.83	0.82	1293
weighted avg	0.86	0.85	0.85	1293

```
In [61]: from sklearn.neighbors import KNeighborsClassifier
KNN= KNeighborsClassifier(n_neighbors=3)
# Train the model using the training sets
KNN.fit(train_x1,train_y1) y_pred1 =
KNN.predict(test_x1)
print("Accuracy Score for KNN : ",
accuracy_score(y_pred1,test_y1)) acc =
accuracy_score(y_pred1,test_y1) pre =
precision_score(test_y1,y_pred1) recall =
recall_score(test_y1,y_pred1) f1 = f1_score(test_y1,y_pred1)
fbeta0_5 = fbeta_score(test_y1,y_pred1,beta=0.5)
fbeta2 = fbeta_score(test_y1,y_pred1,beta=2)
result.loc['KNN_SMOTE'] = [acc,pre,recall,f1,fbeta0_5,fbeta2]
result
```

Accuracy Score for KNN : 0.8371459694989106

Out[61]:

	Accuracy score	Precision	Recall	F1 Score	Fbeta Score(0.5)	Fbeta Score(2)
SVM	0.896365	0.880126	0.744000	0.806358	0.849057	0.767749
SVM_SMOTE	0.944989	0.930208	0.963323	0.946476	0.936648	0.956512
KNN	0.847641	0.709906	0.802667	0.753442	0.726702	0.782225
KNN_SMOTE	0.837146	0.756956	0.997843	0.860866	0.795357	0.938134

```
In [62]: confusion_matrix(test_y1,y_pred1)
```

```
Out[62]: array([[612, 297],
[ 2, 925]], dtype=int64)
```

```
[63]: print(classification_report(test_y1,y_pred1))
```

		precision	recall	f1-score	support
	0	1.00	0.67	0.80	909
1	0.76	1.00	0.86	0.92	927
accuracy				0.84	1836
macro avg	0.88	0.84	0.83		1836
weighted avg	0.88	0.84	0.83		1836

In

Using GridSearchCV

```
In [64]: from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
```

```
In [65]: knn_model=KNeighborsClassifier()
```

```
In [66]: hyperparamters = {'n_neighbors':np.arange(2,10),
                           'p':[1,2]}
rscv_model= RandomizedSearchCV(knn_model,hyperparamters,cv=5)
rscv_model.fit(train_x1,train_y1)
```

```
Out[66]: RandomizedSearchCV(cv=5, estimator=KNeighborsClassifier(),
                             param_distributions={'n_neighbors': array([2, 3, 4, 5, 6, 7,
8, 9]),
                                                  'p': [1, 2]})
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [67]: rscv_model.best_params_
```

```
Out[67]: {'p': 2, 'n_neighbors': 2}
[70]: knn_model= KNeighborsClassifier(n_neighbors=2,p=2)
# Train the model using the training sets
knn_model.fit(train_x,train_y) y_pred =
knn_model.predict(test_x)
print("Accuracy Score for KNN : ",
accuracy_score(y_pred,test_y)) acc =
accuracy_score(y_pred,test_y) pre =
precision_score(test_y,y_pred) recall =
recall_score(test_y,y_pred) f1 = f1_score(test_y,y_pred)
fbeta0_5 = fbeta_score(test_y,y_pred,beta=0.5)
fbeta2 = fbeta_score(test_y,y_pred,beta=2)
result.loc['KNN_SMOTE_Hyperparameter_Tuning'] = [acc,pre,recall,f1,fbeta0_5,fbeta2]
result
```

Accuracy Score for KNN : 0.8592420726991493

```
Out[70]:
```

	Accuracy score	Precision	Recall	F1 Score	Fbeta Score(0.5)	Fbeta Score(2)
SVM	0.896365	0.880126	0.744000	0.806358	0.849057	0.767749
SVM_SMOTE	0.944989	0.930208	0.963323	0.946476	0.936648	0.956512
KNN	0.847641	0.709906	0.802667	0.753442	0.726702	0.782225

In

KNN_SMOTE	0.837146	0.756956	0.997843	0.860866	0.795357	0.938134
KNN_SMOTE_Hyperparameter_Tuning	0.859242	0.793313	0.696000	0.741477	0.771733	0.713505

In

```
[71]: knn_model= KNeighborsClassifier(n_neighbors=2,p=2)
# Train the model using the training sets
knn_model.fit(train_x1,train_y1) y_pred1
= knn_model.predict(test_x1)
print("Accuracy Score for KNN : ",
accuracy_score(y_pred1,test_y1)) acc =
accuracy_score(y_pred1,test_y1) pre =
precision_score(test_y1,y_pred1) recall =
recall_score(test_y1,y_pred1) f1 = f1_score(test_y1,y_pred1)
fbeta0_5 = fbeta_score(test_y1,y_pred1,beta=0.5)
fbeta2 = fbeta_score(test_y1,y_pred1,beta=2)
result.loc['KNN_SMOTE_Hyperparameter_Tuning1'] = [acc,pre,recall,f1,fbeta0_5,fbet
result
```

Accuracy Score for KNN : 0.8932461873638344

Out[71]:

	Accuracy	Precision	Recall	F1 Score	Fbeta	Fbet
	score				Score(0.5)	Score(2)
SVM	0.896365	0.880126	0.744000	0.806358	0.849057	0.76774
SVM_SMOTE	0.944989	0.930208	0.963323	0.946476	0.936648	0.95651
KNN	0.847641	0.709906	0.802667	0.753442	0.726702	0.78222
KNN_SMOTE	0.837146	0.756956	0.997843	0.860866	0.795357	0.93813
KNN_SMOTE_Hyperparameter_Tuning	0.859242	0.793313	0.696000	0.741477	0.771733	0.71350
KNN_SMOTE_Hyperparameter_Tuning1	0.893246	0.829576	0.992449	0.903733	0.857729	0.95495

In []: